

Sentiment Analysis

Learn about sentiment analysis and how it relates to NLP.

Chapter Goals:

- Learn about sentiment analysis and how it's related to NLP
- Create sentiment analysis training pairs from a text corpus

A. Classifying sentiment

Any time you read an opinionated text passage (e.g. social media post, editorial, etc.), the text contains **sentiment**. Sentiment refers to the writer's attitude in the text, so **sentiment analysis** is the task of analyzing or deducing what the writer's attitude is based on the text.

Oftentimes, the writer's attitude can be classified into one of several categories. For example, online user product reviews are normally given on a 1-5 star scale. In this case, we could separate the reviews based on how many stars they have.

An even simpler classification would be to separate the reviews into two categories: positive and negative. This is just regular binary classification, while the stars example is multiclass classification.

“This product is great,
I’ve never found
anything better!”

Positive

“This product is
terrible. It was a
complete waste of my
money.”

Negative

“While this product
isn’t perfect, it suits my
needs well.”

Positive

Classifying product reviews as either "Positive" or "Negative".

B. Training pairs

Similar to regular classification, text classification involves training pairs consisting of input data and labels. In this case, the input data will be tokenized text sequences and each text sequence will be labeled with a class (category). For simplicity, the class labels are just integers in the range $[0, n - 1]$, where n is the total number of classes.

$([1, 5, 6, 8, 2], 0)$

$([3, 5, 2, 9, 8], 1)$

$([2, 4, 7, 6, 6], 0)$

Example batch of 3 training pairs. The maximum sequence length is 5 and there are two classes (binary classification).